

Week 3

I ran the MaxTemperature script on 3 data sets. Following are the outputs and execution time:

```
ubuntu-xenial-large-521 — vagrant@pleasechangeme: ~/hadoop-book/ch02-mr-intro/src/main/java — s
Map input records=5000000
Map output records=4642067
Map output bytes=41778603
Map output materialized bytes=51062785
Input split bytes=904
Combine input records=0
Combine output records=0
Reduce input groups=1
Reduce shuffle bytes=51062785
Reduce input records=4642067
Reduce output records=1
Spilled Records=9284134
Shuffled Maps =8
Failed Shuffles=0
Merged Map outputs=8
GC time elapsed (ms)=3167
CPU time spent (ms)=21880
Physical memory (bytes) snapshot=1924378624
Virtual memory (bytes) snapshot=16911880192
Total committed heap usage (bytes)=1391427584
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
Bytes Read=1030902727
File Output Format Counters
Bytes Written=9

real    1m4.189s
user    0m3.444s
sys     0m0.120s
vagrant@pleasechangeme:~/hadoop-book/ch02-mr-intro/src/main/java$ Hadoop fs -ls /user/$USER/output1
Hadoop: command not found
vagrant@pleasechangeme:~/hadoop-book/ch02-mr-intro/src/main/java$ hadoop fs -ls /user/$USER/output1
-ls: Unknown command
[vagrant@pleasechangeme:~/hadoop-book/ch02-mr-intro/src/main/java$ hadoop fs -ls /user/$USER/output1
Found 2 items
-rw-r--r--  1 vagrant supergroup          0 2017-02-10 20:14 /user/vagrant/output1/_SUCCESS
-rw-r--r--  1 vagrant supergroup          9 2017-02-10 20:14 /user/vagrant/output1/part-r-00000
[vagrant@pleasechangeme:~/hadoop-book/ch02-mr-intro/src/main/java$ hadoop fs -cat /user/$USER/output1/part-r-00000
1990    607
vagrant@pleasechangeme:~/hadoop-book/ch02-mr-intro/src/main/java$
```

```
ubuntu-xenial-large-521 — vagrant@pleasechangeme: ~/hadoop-book/ch02-mr-intro/src
Total time spent by all map tasks (ms)=1816038
Total time spent by all reduce tasks (ms)=318858
Total vcore-seconds taken by all map tasks=1816038
Total vcore-seconds taken by all reduce tasks=318858
Total megabyte-seconds taken by all map tasks=1859622912
Total megabyte-seconds taken by all reduce tasks=326510592
Map-Reduce Framework
  Map input records=35000000
  Map output records=30592765
  Map output bytes=275334885
  Map output materialized bytes=336520775
  Input split bytes=6840
  Combine input records=0
  Combine output records=0
  Reduce input groups=2
  Reduce shuffle bytes=336520775
  Reduce input records=30592765
  Reduce output records=2
  Spilled Records=61185530
  Shuffled Maps =60
  Failed Shuffles=0
  Merged Map outputs=60
  GC time elapsed (ms)=28745
  CPU time spent (ms)=176790
  Physical memory (bytes) snapshot=13301149696
  Virtual memory (bytes) snapshot=114581987328
  Total committed heap usage (bytes)=9769713664
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=7993006187
File Output Format Counters
  Bytes Written=18

real    6m40.488s
user    0m4.712s
sys     0m0.288s
[vagrant@pleasechangeme:~/hadoop-book/ch02-mr-intro/src/main/java$ hadoop fs -cat /user/$USER/output3/par
1990    607
1992    605
vagrant@pleasechangeme:~/hadoop-book/ch02-mr-intro/src/main/java$
```

Retired Jobs

Show 20 ▾ entries										Search: <input type="text"/>	
Submit Time ▾	Start Time ▾	Finish Time ▾	Job ID ▾	Name ▾	User ▾	Queue ▾	State ▾	Maps Total ▾	Maps Completed ▾	Reduces Total ▾	Reduces Completed ▾
2017.02.10 21:29:11 CST	2017.02.10 21:29:19 CST	2017.02.10 21:40:12 CST	job_1486782805443_0001	Max temperature	vagrant	default	SUCCEEDED	88	88	1	1
Submit Time	Start Time	Finish Time	Job ID	Name	User	Queue	State	Maps T	Maps Comple	Reduces Tc	Reduces Com
Showing 1 to 1 of 1 entries											First Previous 1 Next Last